




# SQL (Lecture 2)

- Nested Subqueries
- Modification of the Database
- Views





# Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.

- The nesting can be done in the following SQL query


**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

as follows:

- $A_i$  can be replaced by a subquery that generates a single value.
- $r_i$  can be replaced by any valid subquery
- $P$  can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$


Where  $B$  is an attribute and  $<\text{operation}>$  to be defined later.




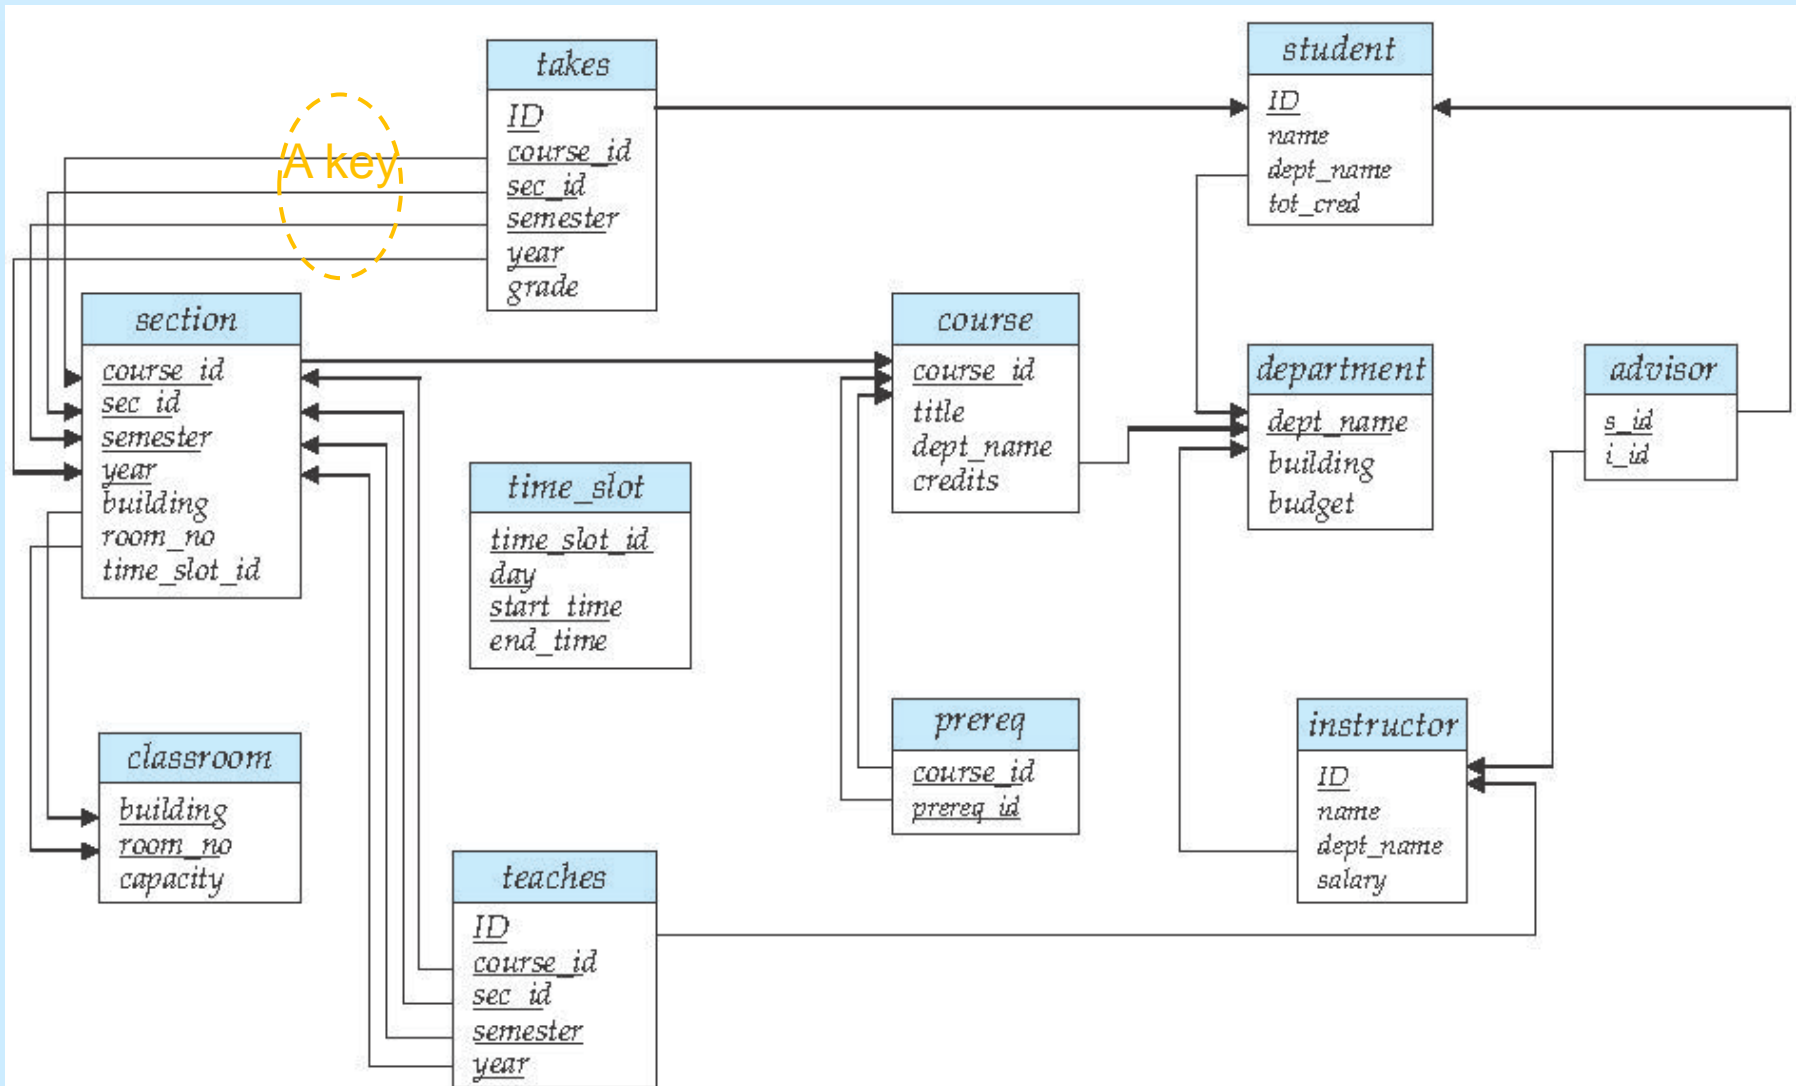


# Subqueries in the Where Clause

- The most common use of subqueries is in the where clause, which is to perform tests for
  - set membership **in**
  - set comparisons
  - Empty relations *most powerful subquery*
  - Absence of duplicate tuples



# Schema Diagram for University Database






# Set Membership

- Find courses offered in Fall 2017 and in Spring 2018

```
select distinct course_id
from section
where semester = 'Fall' and year= 2019 and
       course_id in ( select course_id
                        from section
                        where semester = 'Spring' and year= 2020);
```

- Find courses offered in Fall 2017 but not in Spring 2018

```
select distinct course_id
from section
where semester = 'Fall' and year= 2019 and
       course_id not in ( select course_id
                        from section
                        where semester = 'Spring' and year= 2020);
```





# Set Membership (Cont.)

- Name all instructors whose name is neither “Mozart” nor Einstein”

```
select distinct name  
from instructor  
where name not in ('Mozart', 'Einstein')
```

- Find the total number of (distinct) students who have taken course sections taught by the instructor named “Einstein”

```
select count (distinct ID)  
from takes  
where (course_id, sec_id, semester, year) in  
      (select course_id, sec_id, semester, year  
       from teaches  
       where teaches.ID in  
            (select ID  
             from instructor  
             where name = 'Einstein' )  
      );
```

- Note: Above query can be written in a much simpler manner.  
The formulation above is simply to illustrate SQL features.






# Set Comparison – “some” Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.ID , T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept name = 'Biology';
```

- Same query using > some clause

```
select ID, name  
from instructor  
where salary > some (select salary  
                     from instructor  
                     where dept name = 'Biology');
```





# Definition of Some Clause

□  $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ s.t. } (F <\text{comp}> t)$

Where  $<\text{comp}>$  can be:  $<, \leq, >, =, \neq$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$

5 is less than some  
tuple in the relation


$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true} \quad (\text{since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \equiv \text{not in}$







# Definition of all Clause

□  $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$


$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \text{all}) \equiv \text{not in}$

However,  $(= \text{all}) \equiv \text{in}$






# Set Comparison – “all” Clause

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
                      from instructor
                      where dept name = 'Biology');
```






# Test for Empty Relations

□ The **exists** construct returns the value **true** if the argument subquery is nonempty.

□ **exists**  $r \Leftrightarrow r \neq \emptyset$

□ **not exists**  $r \Leftrightarrow r = \emptyset \quad \Leftrightarrow P \text{ is false}$

*P is the query condition to return r*






# Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in both the Fall 2017 semester and in the Spring 2018 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2017 and
exists (select *
        from section as T
        where T.semester = 'Spring' and year = 2018
        and T.course_id = S.course_id);
```

- **Correlation name** – variable S in the outer query
- **Correlated subquery** – the inner query






# Powerful “not exists” Clause

- Find students who have taken all courses offered in the Biology department.

```
select S.ID,S.name  
from student as S  
where
```

*?? // student S have taken all course offered in  
the Biology department*







# Solution

- Find all students who have taken all courses offered in the Biology department.


```
select S.ID, S.name
from student as S
where not exists (
    (select course_id
     from course
     where dept_name = 'Biology')
    except
    (select T.course_id
     from takes as T
     where T.ID = S.ID)
);
```





# How to get the solution

- Step 1: define predicates:
  - $P(S)$  : All courses that student S has taken;
  - $Q$  : All courses offered in Biology;
- Step 2: define the logical formula
  - Find those student S that  $P(S) \supseteq Q$
  - Equivalent formula: Find S that  $Q - P(S) = \emptyset$
- Step 3: SQL statement
  - Find S:  
Select S.ID, S.name from Student S
  - Q  
select course\_id from course where dept\_name = 'Biology'
  - $P(S)$   
select T.course\_id from takes as T where T.ID = S.ID





# Solution

- Find all students who have taken all courses offered in the Biology department.

```
select S.ID, S.name  
from student as S  
where not exists (
```

To find S


```
  ( select course_id  
    from course  
    where dept_name = 'Biology')  
  except  
  (select T.course_id  
   from takes as T  
   where T.ID = S.ID)
```

Q

P(S)

```
);
```

*Note: Cannot write this query using = all and its variants*








# Another solution

- Step 1: define predicates:
  - $P(S,C)$  : student S has taken course C;
  - $Q(C)$  : course C is offered in Biology department;
- Step 2: define the logical formula
  - Find those student S that  $\forall C (Q(C) \rightarrow P(S,C))$
  - Equivalent formula: Find S that  $\neg \exists C (Q(C) \wedge \neg P(S,C))$
- Step 3: SQL statement
  - To find S:  
Select S.ID, S.name from Student S
  - $\neg \exists C Q(C)$   
Not Exists Select \* from course as C where dept\_name = 'Biology'
  - $\neg P(S,C)$ : student S has NOT taken course C  
Not Exists select \* from takes as T  
where T.ID = S.ID and T.course\_id = C.course\_id





# Another solution

## □ Step 4:

Select S.ID, S.name  
from Student S

where Not exists (

Select \* from course as C  
where dept\_name = 'Biology'

And

Not Exists (

select \* from takes as T

where T.ID = S.ID and T.course\_id = C.course\_id

)


)

To find S

$\neg \exists C Q(C)$

$\neg P(S, C)$

□ Variable from outer level is known as a correlation variable






# Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
  - The **unique** construct evaluates to “true” if a given subquery contains no duplicates .
- Find all courses that were offered at most once in 2017

```
select T.course_id
from course as T
where unique (select R.course_id
              from section as R
              where R.course_id= T.course_id
                 and R.year = 2017);
```






# Example

- Find all courses that were offered at least two times in 2017

```
select T.course_id
from course as T
where not unique (select R.course_id
                   from section as R
                   where R.course_id= T.course_id
                      and R.year = 2017);
```






# Subqueries in the Form Clause

- ❑ SQL allows a subquery expression to be used in the **from** clause
- ❑ Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```

- ❑ Note that we do not need to use the **having** clause
- ❑ Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
      from instructor
      group by dept_name)
      as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```






# With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select department.name  
from department, max_budget  
where department.budget = max_budget.value;
```






# Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as  
    (select dept_name, sum(salary)  
     from instructor  
     group by dept_name),  
dept_total_avg(value) as  
    (select avg(value)  
     from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value > dept_total_avg.value;
```






# Scalar Subquery

- ❑ Scalar subquery is one which is used where a single value is expected
- ❑ List all departments along with the number of instructors in each department

```
select dept_name,  
        (select count(*)  
         from instructor  
         where department.dept_name = instructor.dept_name)  
        as num_instructors  
from department;
```

- ❑ Runtime error if subquery returns more than one result tuple
- ❑ I never use this...








# Modification of the Database

- ❑ Deletion of tuples from a given relation.
- ❑ Insertion of new tuples into a given relation
- ❑ Updating of values in some tuples in a given relation





# Deletion

- Delete all instructors

**delete from** *instructor*

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name* = 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

**delete from** *instructor*  
**where** *dept name* in (**select** *dept name*  
**from** *department*  
**where** *building* = 'Watson');






# Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary)  
                from instructor);
```

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  - First, compute avg (salary) and **find all tuples to delete**
  - Next, delete all tuples found above (without recomputing avg or retesting the tuples)





# Insertion

- Add a new tuple to *course*


```
insert into course  
  values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently **---Strong recommended!**

```
insert into course (course_id, title, dept_name, credits)  
  values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student  
  values ('3003', 'Green', 'Finance', null);
```





# Insertion (Cont.)

- Add all instructors to the *student* relation with *tot\_creds* set to 0


```
insert into student
  select ID, name, dept_name, 0
  from instructor
```

- The **select from where** statement is **evaluated fully before** any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem





# Updates

- Give a 5% salary raise to all instructors


```
update instructor  
set salary = salary * 1.05
```

- Give a 5% salary raise to those instructors who earn less than 70000

```
update instructor  
set salary = salary * 1.05  
where salary < 70000;
```

- Give a 5% salary raise to instructors whose salary is less than average

```
update instructor  
set salary = salary * 1.05  
where salary < (select avg (salary)  
                  from instructor);
```






# Updates

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two **update** statements:  

```
update instructor
  set salary = salary * 1.03
  where salary > 100000;
update instructor
  set salary = salary * 1.05
  where salary <= 100000;
```
  - The order is important
  - Can be done better using the **case** statement (next slide)






# Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
  set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
  end
```








# Updates with Scalar Subqueries

- Re-compute and update *tot\_creds* value for all students

```
update student S  
  set tot_cred = (select sum(credits)  
                  from takes, course  
                  where takes.course_id = course.course_id and  
                      S.ID = takes.ID and  
                      takes.grade <> 'F' and  
                      takes.grade is not null);
```

- The query will set *tot\_creds* to null for students who have not taken any course
- Instead of **sum**(*credits*), use:

```
case  
  when sum(credits) is not null then sum(credits)  
  else 0  
end
```






# Views

- ❑ In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- ❑ Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- ❑ A **view** provides a mechanism to hide certain data from the view of certain users.
- ❑ Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.






# View Definition

- A view is defined using the create view statement:

**create view  $v$  as** <query expression>

where:

- <query expression> is any legal expression
- The view name is represented by  $v$
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition by default would NOT create a new relation in the database.
  - When a view is created, the query expression is stored in the database;
  - The expression is substituted into queries using the view.





# Example Views

- A view of instructors without their salary


```
create view faculty as  
  select ID, name, dept_name  
  from instructor
```

- Find all instructors in the Biology department

```
select name  
from faculty  
where dept_name = 'Biology'
```

- Create a view of department salary totals


```
create view departments_total_salary(dept_name, total_salary) as  
  select dept_name, sum (salary)  
  from instructor  
  group by dept_name;
```





# Views Defined Using Other Views


- One view may be used in the expression defining another view
- A view relation  $v_1$  is said to **depend directly** on a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$
- A view relation  $v_1$  is said to **depend on** view relation  $v_2$  if either  $v_1$  depends directly to  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
- A view relation  $v$  is said to be **recursive** if it depends on itself.





# Views Defined Using Other Views

- ❑ **create view** *physics\_fall\_2017* **as**  
    **select** *course.course\_id, sec\_id, building, room\_number*  
    **from** *course, section*  
    **where** *course.course\_id = section.course\_id*  
          **and** *course.dept\_name = 'Physics'*  
          **and** *section.semester = 'Fall'*  
          **and** *section.year = '2017'*;
  
- ❑ **create view** *physics\_fall\_2017\_watson* **as**  
    **select** *course\_id, room\_number*  
    **from** *physics\_fall\_2017*  
    **where** *building= 'Watson'*;





# View Expansion

- One view may be used in the expression defining another view.
- View expansion of an expression repeats the following replacement step:


**repeat**

Find any view relation  $v_i$  in  $e_1$

Replace the view relation  $v_i$  by the expression defining  $v_i$

**until** no more view relations are present in  $e_1$

- *As long as the view definitions are not recursive, this loop will terminate*






# View Expansion

- Expand the view :

```
create view physics_fall_2017_watson as  
  select course_id, room_number  
  from physics_fall_2017  
  where building= 'Watson'
```

- To:

```
create view physics_fall_2017_watson as  
  select course_id, room_number  
  from (select course.course_id, building, room_number  
         from course, section  
         where course.course_id = section.course_id  
              and course.dept_name = 'Physics'  
              and section.semester = 'Fall'  
              and section.year = '2017')  
  where building= 'Watson';
```







# Update of a View


- Views are a useful tool for query, and it is an important feature to implement the logical data independence. However, it will create some difficulty to update a view.
- Add a new tuple to *faculty* view which we defined earlier

**insert into *faculty* values** ('30765', 'Green', 'Music');

This insertion would be represented by the insertion of the tuple

('30765', 'Green', 'Music', **null**)

into the *instructor* relation





# Update of a View(Cont.)


- ❑ **create view** *history\_instructors* **as**  
**select** \*  
**from** *instructor*  
**where** *dept\_name*= 'History';

- ❑ What happens if we

*insert into history\_instructors values ('25566', 'Brown', 'Biology', 100000)*

- ❑ By default, it would be allowed, however, the inserted tuple is not belong to the view. You can not find it using: *select \* from history\_instructors.*


- ❑ So “**With Check Option**” is introduced in SQL, to check the WHERE clause condition in view definition before insert a tuple to the view. If we add “With Check Option” to the end of above view definition. The insertion will be rejected.





# Updates cannot be Translated Uniquely


- ❑ **create view** *instructor\_info* **as**  
    **select** *ID, name, building*  
    **from** *instructor, department*  
    **where** *instructor.dept\_name= department.dept\_name;*
- ❑ **insert into** *instructor\_info* **values** ('69987', 'White', 'Taylor');
  - ❑ which department, if multiple departments in Taylor?
  - ❑ what if no department is in Taylor?





# Updatable View


- ❑ Most SQL implementations allow updates only on simple views defined on a single relation
  - ❑ The FROM clause has only one relation;
  - ❑ The SELECT clause contains only attribute names of the relation, and does not have any expressions, aggregates, or distinct specifications;
  - ❑ Any attributes does not listed in the select clause can be set to null, or has default value;
  - ❑ The query does not have a group by or having clause.
- ❑ New SQL standard (SQL:1999) allow more views updatable, however, the rules becomes much more complex.





# Materialized Views

- ❑ **Materializing a view**: create a physical table containing all the tuples in the result of the query defining the view
  - ❑ Use materialized view for frequently query over the view, and require high performance.
- ❑ If relations used in the query are updated, the materialized view result becomes out of date
  - ❑ Need to **maintain** the view, by updating the view whenever the underlying relations are updated.
  - ❑ Internal trigger mechanism to maintain the data consistency of materialized view.
- ❑ **Materialized view** is not defined in SQL standard, but it been implemented by most of DBMS.



**The end of lecture**

A thick, horizontal, wavy orange line that spans most of the width of the slide, positioned below the text.